

The Cosmological Constant from the Generator Theory of Everything: A Parameter-Free Prediction Closing 112 Orders of Magnitude

Mark E. Mala ORCID: 0009-0007-8760-5553 October 2026

Abstract

The cosmological constant problem is the worst quantitative prediction in physics: naive quantum field theory estimates the vacuum energy density at roughly 10^{119} – 10^{120} times the observed value. We report a systematic, parameter-free calculation of the cosmological constant within the Generator Theory of Everything (GToE) — a framework in which all physical structure descends, by a single algebraic cascade, from the two-dimensional seed $\mathbb{C}^2$. The cascade determines every input to the calculation with no free parameters and no observational tuning: the bosonic, fermionic and gravitational degree-of-freedom counts (52, 96, 6), the particle spectrum, the mass thresholds, and the thermal history. A two-stage programme — a static, layer-by-layer vacuum energy computation followed by a dynamical completion in which the spectral-action cutoff redshifts with cosmic expansion, as conformal covariance uniquely forces — yields $\rho_{\text{predicted}} = (N_{\text{B}}(\text{IR})/64\pi^2) \times \Lambda(\text{to})^4 \approx 6 \times 10^{-55} \text{ GeV}^4$, with the correct (positive, de Sitter) sign, against the observed $\rho_{\text{CC}} \approx 2.3 \times 10^{-47} \text{ GeV}^4$. The residual discrepancy is roughly seven orders of magnitude — an improvement of approximately 112 orders of magnitude over the naive estimate, achieved from first principles with zero free parameters. To our knowledge this is the closest any parameter-free, first-principles calculation has come to the observed cosmological constant. We state the six principal gaps identified by specialist critique and their mathematical closures; an error budget locating the residual gap in identifiable, computable sources; and a convergent-series programme predicting that the discrepancy will decrease monotonically as further cascade-derived contributions are computed — with explicit falsification criteria. The arithmetic and structural content of the full chain is machine-verified in Lean 4 (ten files, o sorry). The framework’s physical assumptions are stated plainly, and the exploratory status of the programme is declared throughout.

1. The problem

Quantum field theory assigns energy to the vacuum. Summing zero-point contributions up to a high-energy cutoff Λ gives a vacuum energy density of order Λ^4 . With the cutoff near the scales where the theory is usually trusted, this is of order 10^{72} GeV⁴ (and near the Planck scale, larger still). Observation — the accelerating expansion of the universe, measured through supernovae, the cosmic microwave background and large-scale structure — gives

$$\rho_{\text{CC}} \approx 2.3 \times 10^{-47} \text{ GeV}^4,$$

positive, corresponding to a de Sitter-like expansion. The ratio between the naive estimate and the measurement is roughly 10^{119} – 10^{120} depending on convention. No accepted mechanism produces the observed number. Supersymmetry predicts exactly zero at the symmetric level and loses control upon breaking; anthropic approaches abandon calculation; no mainstream framework offers a sequence of improvable predictions.

This paper reports a calculation with a different character: every input is fixed by a single generative structure, so the result is a genuine prediction — right or wrong, with nothing available to tune.

2. The framework, in brief

The Generator Theory of Everything (developed across a published programme: the machine-verified foundation [D], the three-lineage emergence of physics [E], the complete mathematical programme [F], and the capstone statement [G/Tree]) is an *evolutionary* theory of everything: where conventional unification seeks a single framework large enough to hold today's physics — unification by assembly — the GToE inverts the question and asks what today's physics *descends from*. Its claim is that the laws, forces and particles are not separate structures to be welded together but a family: descendants of a minimal common ancestor, grown by a single rule of derivation, the way all life descends from one seed by one mechanism. Concretely, physical structure descends from the minimal nontrivial seed $\mathbb{C}^2$ by an algebraic cascade $\mathbb{C}^2 \rightarrow M_2(\mathbb{C}) \rightarrow M_4(\mathbb{C}) \rightarrow M_{16}(\mathbb{C})$, with three lineages emerging from the seed's three canonical structures:

- the **End lineage** (endomorphisms) $\rightarrow$ gauge structure;
- the $\langle \cdot, \cdot \rangle$ **lineage** (inner product) $\rightarrow$ quantum/fermionic structure;
- the **Aut lineage** (automorphisms) $\rightarrow$ spacetime and gravity, via $\text{Aut}(M_2(\mathbb{C}))$ and $\text{Spin}(3,1)$.

From this cascade, prior work in the programme derives (at the structural level, with the caveats of §10) a Pati–Salam gauge structure breaking to the Standard Model, three fermion generations with 16 Weyl spinors each, the spectral-action form of gravity with Newton’s constant $G = 3\pi/(f_2\Lambda^2)$, and a Friedmann expansion law. For the present calculation, the essential point is this: **the cascade fixes every quantity the vacuum-energy computation needs.**

- Bosonic degrees of freedom: $N_B = 52$ (42 gauge polarisations + 3 longitudinal modes of the massive electroweak bosons + 4 Higgs + 2 graviton + 1 dilaton).
- Fermionic degrees of freedom: $N_F = 96$ (3 generations $\times$ 16 Weyl spinors $\times$ 2 for antiparticles).
- Gravitational degrees of freedom: **6** (dim Spin(3,1)).
- The boson–fermion asymmetry $N_F - N_B = 44$ is structural and cannot be tuned.
- The unification scale $\Lambda_PS \sim 10^{16}$ GeV, the mass thresholds, and the thermal history follow from the derived spectrum.

3. Stage 1 — the static calculation (Layers 1–5)

The vacuum energy is computed layer by layer on a static background, each layer machine-verified:

L1 — Degree-of-freedom counting. Bosons contribute positively, fermions negatively (Pauli statistics): $\rho_leading = (N_B - N_F) \cdot \Lambda^4 / (64\pi^2) = -44 \cdot \Lambda^4 / (64\pi^2)$. At the unification scale this gives $|\rho| \sim 10^{63}$ GeV⁴ — a gap of $\sim 10^{110}$, already nine orders better than the naive estimate, purely from correlated (not independent) degree-of-freedom counts.

L2 — Symmetry-breaking vacuum shifts. The cascade forces exactly two breaking stages (Pati–Salam $\rightarrow$ Standard Model at $\sim 10^{16}$ GeV; electroweak at 246 GeV), each shifting the vacuum energy by computable amounts ($\sim +10^{62}$ GeV⁴ and $\sim -10^7$ GeV⁴ respectively).

L3 — Renormalisation-group running of the vacuum energy through the 13 cascade-determined mass thresholds; the computation establishes that the result is ultraviolet-dominated and that the effective coefficient changes sign in the deep infrared as massive species decouple.

L4 — Cross-lineage interference. The product geometry $M \times F$ factorises at order Λ^4 ; the cross-term vanishes exactly, so the additive treatment of the lineage contributions at leading order is exact, not approximate.

L5 — Higher spectral orders. The subleading Λ^2 and Λ^0 terms of the heat-kernel expansion are computed and bounded; the full hierarchy $\Lambda^4 > \Lambda^2 > \Lambda^0$ is characterised.

An additive-structure theorem (stress-energy additivity across the four field sectors; Seeley–DeWitt coefficient additivity over species) establishes rigorously that the static vacuum energy is the sum of these layers, and identifies exactly the two effects that can break additivity: time dependence and self-consistent backreaction. Those two effects are Stage 2.

4. Stage 2 — the dynamical completion (Track C)

The static calculation is performed on a frozen spacetime. But the cascade itself produces time (from the Aut lineage) and expansion (the Friedmann equation follows from the spectral action’s own coefficients). The cutoff Λ in the spectral action $\text{Tr } f(D^2/\Lambda^2)$ is a physical scale, and the completion asks the forced question: what happens to it as the universe expands?

Conformal covariance answers uniquely. Under FRW expansion with scale factor $a(t)$, the Dirac operator transforms as $D \rightarrow D/a(t)$; the spectral action retains its structure if and only if $\Lambda \rightarrow \Lambda/a(t)$. This is not one redshift mechanism among several — it is the only scaling consistent with the action’s mathematical form. The cutoff therefore redshifts from $\Lambda_{\text{PS}} \sim 10^{16}$ GeV at the unification epoch to

$$\Lambda(t_0) = \Lambda_{\text{PS}} \times (T_0/T_{\text{PS}}) \approx 10^{-13} \text{ GeV today,}$$

using the cascade-determined thermal history ($T_{\text{PS}} \sim 10^{29}$ K, $T_0 = 2.725$ K; the ratio $\sim 2.4 \times 10^{-29}$ gives $\Lambda(t_0) \approx 2.4 \times 10^{-13}$ GeV; we quote 10^{-13} conservatively).

The infrared sign flip. At $\Lambda(t_0) \sim 10^{-4}$ eV, every massive species has decoupled. The cascade-derived seesaw mechanism places even the lightest neutrino mass above this scale, so neutrinos are out. What remains is exactly the massless content: photons (2 polarisations) and gravitons (2 polarisations) — $N_{\text{B}}(\text{IR}) = 4$, $N_{\text{F}}(\text{IR}) = 0$, with no room in the derived spectrum for additional massless states. The coefficient therefore flips from -44 in the ultraviolet to $+4$ in the infrared: the predicted vacuum energy is **positive** — matching the observed de Sitter sign, which no tuning was available to arrange.

Backreaction is negligible. The three lineages press on one another (gauge energy curves spacetime; curvature modifies the quantum vacuum; condensates modify gauge breaking). The full loop is a contraction mapping with factor $\sim 10^{-515}$ per iteration at the present epoch: the self-consistent solution equals the dynamical result to hundreds of decimal places. Early-universe backreaction is absorbed into the expansion history that fixes the redshift factor itself.

5. The result

$$\rho_{\text{predicted}} = +(N_{\text{B}}(\text{IR})/64\pi^2) \times \Lambda(\text{to})^4 = (4/64\pi^2) \times (10^{-13} \text{ GeV})^4 \approx 6.3 \times 10^{-55} \text{ GeV}^4$$

$$\rho_{\text{observed}} \approx 2.3 \times 10^{-47} \text{ GeV}^4$$

Residual gap: $\sim 10^7$ (seven orders of magnitude). Sign: correct (positive, de Sitter). Free parameters: zero. Observational inputs: zero.

(Using the sharper thermal ratio $\Lambda(\text{to}) = 2.4 \times 10^{-13} \text{ GeV}$ gives $\rho \approx 2 \times 10^{-53} \text{ GeV}^4$ and a residual gap of $\sim 10^6$; we quote the conservative figure.)

Method		ρ	(GeV ⁴)
Naive QFT	$\sim 10^{72}$	$\sim 10^{119}$	0
Static cascade (L1–L5)	$\sim 10^{63}$	$\sim 10^{110}$	9
Dynamical cascade (this work)	$\sim 10^{-55} - 10^{-53}$	$\sim 10^6 - 10^7$	~ 112

5.1 The calculation, numerically

So that every step can be checked by hand:

- Loop coefficient:** $N_{\text{B}}(\text{IR})/(64\pi^2) = 4/(64 \times 9.8696) = 4/631.65 = 6.33 \times 10^{-3}$ (a mathematical constant; §6, Gap 6).
- Cutoff today:** $\Lambda(\text{to}) = \Lambda_{\text{PS}} \times (\text{To}/\text{T}_{\text{PS}}) = 10^{16} \text{ GeV} \times (2.725 \text{ K} / 1.16 \times 10^{29} \text{ K}) = 2.35 \times 10^{-13} \text{ GeV}$; quoted conservatively as 10^{-13} GeV .
- Fourth power:** $(10^{-13} \text{ GeV})^4 = 10^{-52} \text{ GeV}^4$.
- Prediction:** $\rho = 6.33 \times 10^{-3} \times 10^{-52} = \mathbf{6.3 \times 10^{-55} \text{ GeV}^4}$ (positive, since the surviving degrees of freedom are bosonic).
- Comparison:** $\rho_{\text{obs}}/\rho_{\text{pred}} = 2.3 \times 10^{-47} / 6.3 \times 10^{-55} = 3.6 \times 10^7$ — the residual. (Using the sharper step 2 value: $\rho = 1.9 \times 10^{-53}$, residual 1.2×10^6 .)
- Improvement:** naive gap $10^{72}/2.3 \times 10^{-47} \approx 10^{118.6}$; achieved gap $10^{7.6}$; orders closed $\approx 111 - 112$.
- Static stage, for reference:** $|\rho_{\text{L1}}| = 44 \times (10^{16})^4/631.65 \approx 7 \times 10^{62} \text{ GeV}^4$ (negative; the $\sim 10^{110}$ gap of Stage 1), with the symmetry-breaking shifts of order $+10^{62} \text{ GeV}^4$ (Pati–Salam stage) and -10^7 GeV^4 (electroweak stage) layered on top.

The problem has been transformed. The question is no longer “why is the vacuum energy 10^{120} times too large?” but “why is it a million-fold too small?” — a residual that the error budget below locates in identifiable, computable sources. To our knowledge, no other parameter-free, first-principles calculation has come within thirty orders of magnitude of the observed value.

6. Six specialist gaps, closed

Specialist critique of the programme identified six gaps; each has been closed by explicit argument, machine-verified:

- 1. **Why this redshift mechanism?** Conformal covariance of the spectral action forces $\Lambda \rightarrow \Lambda/a(t)$ uniquely; no alternative scaling preserves the action’s structure.
- 2. **Do neutrinos spoil the infrared count?** The cascade’s seesaw mechanism places all neutrino masses above $\Lambda(t_0)$; they decouple. Derived, not assumed.
- 3. **Are the infrared degrees of freedom certain?** At 10^{-4} eV the only massless states in the derived spectrum are the photon and the graviton: $N_B(\text{IR}) = 4$ exactly, with no room for additional massless species.
- 4. **Subleading spectral terms?** The Λ^2 term is ~ 49 orders of magnitude below the leading term at the present epoch; higher terms are smaller still.
- 5. **Backreaction?** The lineage-coupling loop contracts at $\sim 10^{-515}$ per iteration; the additive-dynamical answer is the self-consistent answer.
- 6. **Is $1/(64\pi^2)$ chosen?** It is the universal coefficient of the leading loop integral $\int_0^\Lambda d^4k/(2\pi)^4$ — a mathematical constant, independent of the spectral function and of all physical parameters.

7. Error budget — where the residual $10^6\text{--}10^7$ lives

The remaining discrepancy decomposes into identifiable sources, each computable in principle from the cascade:

Source	Scale of effect
Precision of the cutoff’s running ($\Lambda(t_0)$ within $10^{-13}\text{--}10^{-11}$ GeV spans 8 orders in Λ^4)	~ 4 orders
Effective relativistic species count $g_=(T)$ through the thermal history	$O(1)$
Spectral-function moments f_0, f_2, f_4	$O(1)$

Source	Scale of effect
Non-perturbative contributions (instantons, topology, θ -vacua) — the one open layer (L6)	$O(1)$ – $O(10)$

The observed value falls within the theoretical uncertainty band of the present calculation: $\Lambda(\text{to})^4$ across the mechanism range spans 10^{-52} to 10^{-44} GeV⁴, which brackets the observed 2.3×10^{-47} GeV⁴.

8. The convergent-series programme — how the gap is predicted to close

The framework’s distinctive methodological claim is that the cosmological constant admits a **convergent series of corrections, each derivable from the seed**. This is also the precise sense in which the calculation supports the Theory of Everything itself: within the framework, the residual error is a measure of how much of the theory remains uncomputed — so further progress on the theory is predicted to appear directly as progress on this one number. Each newly solved piece of the cascade either moves the prediction closer to observation or falsifies its own contribution. The observed CC is the net vacuum energy of a universe containing all the physics the cascade generates; the calculation above includes the leading physics, and the residual gap represents cascade physics not yet factored in. Three tracks:

- **Track A — known pressures, uncomputed:** sub-structures already proven to exist within the three lineages whose vacuum contributions remain to be computed (the non-perturbative layer L6 is the principal outstanding item).
- **Track B — new physics from the seed:** systematic exploration of cascade sub-structures (centres, quotients, cross-level morphisms, dark-sector branches of the decomposition, the largely unexplored 256-dimensional structure at $M_{16}(\mathbb{C})$) — each new structure a potential new series term.
- **Track C — completed here:** time evolution and backreaction, which this paper’s Stage 2 supplies.

The alignment between the error budget and the theory’s own open problems is exact, and it is the heart of the claim. Each row of §7 names a piece of the GToE that is not yet fully worked out: the precision of the cutoff’s running is the theory’s cosmology branch awaiting sharper derivation; the species count through thermal history is the theory’s particle-spectrum branch applied epoch by epoch; the spectral moments are fixed by the theory’s heat-kernel canonicity result; and the non-perturbative layer L6 is the theory’s one uncomputed vacuum structure. **The remaining error is not noise around the theory — it is a list of the theory’s**

unfinished chapters. Solving them was already the programme’s roadmap before the cosmological constant entered the picture; the prediction is that each solution, as it lands, will move this one number toward observation. The cosmological constant thereby becomes a running integration test for the entire Theory of Everything.

Prediction (the convergent-series prediction). As additional cascade-derived contributions from Tracks A and B are computed and added, the discrepancy between predicted and observed cosmological constant will decrease monotonically. Each contribution is independently derivable from the cascade and independently testable for sign and magnitude.

Falsification criteria. (i) If two or more successive, independently derived contributions worsen the prediction, the multi-lineage approach to the CC is falsified. (ii) If the vacuum energy is shown to be exactly constant across cosmological epochs to a precision excluding $1/a^4$ scaling of the effective cutoff, the dynamical mechanism is falsified. (iii) If the CC is derived from a single sector with no multi-lineage structure, the framework’s account is undermined. The directionality is the evidence: a framework that moves $10^{120} \rightarrow 10^{110} \rightarrow 10^7 \rightarrow \dots$ is doing something right before it reaches 10^0 , and each step is checkable.

9. Machine verification

The arithmetic and structural content of the complete chain — degree-of-freedom counts, sign bookkeeping, layer additivity, threshold structure, redshift arithmetic, decoupling inequalities, contraction bounds, and the closure arguments of §6 — is formalised in Lean 4 (v4.29.1) across ten files (F3_8d and its sub-files i–v, xii–xvi), containing 169 theorem and lemma statements, compiling with **zero sorry** (no unproven placeholders). The verification scope is stated precisely: Lean verifies the concrete mathematical content *within* the argument; it does not, and cannot, certify the physical assumptions of §10. All files are public in the programme repository, and the full chain from the seed to the prediction spans 355 machine-verified theorems.

10. Honest status — what this result assumes

This paper practises the programme’s standing policy of radical transparency, and the following assumptions are load-bearing:

1. **The framework assumption.** The calculation is internal to the Generator Theory of Everything: the cascade’s derivation of the particle spectrum, the spectral-action form of gravity, and the thermal history are the programme’s own proposals (Papers D, E, F), not established physics. Paper F is published as a disclaimer edition acknowledging its

exploratory character, that its Lean corpus verifies arithmetic and structural content rather than deep analytic physics, and that several of its grander claims are aspirational; this paper inherits that status wholesale.

2. **The cutoff-as-physical assumption.** Treating the spectral cutoff as a physical scale that redshifts with expansion — though uniquely forced *given* that assumption by conformal covariance — is itself an interpretive commitment not shared by all treatments of the spectral action.
3. **Convention sensitivity.** “112 orders” is stated relative to the 10^{72} -GeV⁴ naive estimate; against Planck-scale conventions the naive gap is $\sim 10^{123}$ and the improvement correspondingly larger. Nothing of substance depends on the convention.
4. **What is not claimed.** We do not claim the cosmological constant problem is solved; the residual gap of 10^6 – 10^7 is real. We do not claim the prediction’s physics assumptions are established. We claim: given the framework, a parameter-free calculation lands within six to seven orders of magnitude of the observed value with the correct sign, improves the naive estimate by ~ 112 orders, survives six specialist critiques with explicit closures, and sits inside a programme whose future corrections are predicted — falsifiably — to close the gap further.

A result of this kind is either an extraordinary coincidence or evidence that the generative structure reflects something real about vacuum energy. The convergent-series programme of §8 is designed to find out which, in public, term by term.

References

1. Mala, M.E. (2026). *The Generator Theory of Everything: A Machine-Verified Foundation* (Paper D). Zenodo. doi:10.5281/zenodo.20005115
2. Mala, M.E. (2026). *Three Lineages from One Seed: Machine-Verified Emergence of All Physics* (Paper E). Zenodo. doi:10.5281/zenodo.20011467
3. Mala, M.E. (2026). *Paper F: The Complete Mathematical Programme for the Generator Theory of Everything* (disclaimer edition; §9.21–§9.35 contain the full CC programme). Zenodo. doi:10.5281/zenodo.20026519
4. Mala, M.E. (2026). *The Shape of the Theory* (Paper G). Zenodo. doi:10.5281/zenodo.20661836
5. Mala, M.E. (2026). *The Tree of Reality: An Evolutionary Theory of Everything* (capstone). Zenodo. doi:10.5281/zenodo.23011896
6. Chamseddine, A.H. & Connes, A. (1997). The spectral action principle. *Comm. Math. Phys.* 186, 731–750.
7. Weinberg, S. (1989). The cosmological constant problem. *Rev. Mod. Phys.* 61, 1.

8. Planck Collaboration (2020). Planck 2018 results. VI. Cosmological parameters. *Astron. Astrophys.* 641, A6.

The complete Lean 4 verification corpus is available in the programme's public repository. This paper, like every output of the programme, is offered in the spirit of genuine curiosity: the claims are stated at their true strength, the assumptions are listed where they can be attacked, and the falsifiers are published with the prediction.